

# Sentiment Identification

*Practical 9: A Sentimental Journey*

# Human Ratings

## Q1

# Do a Rating Task

- ◆ Get 3 classmates (opinion holders) to write three different opinions about their phone
- ◆ Get 3 diff people (raters) to rate these comments as positive, negative, neutral or can't-say
- ◆ Take this  $3 \times 3$  matrix and find the inter-rater reliability between your 3 raters using Kappa
- ◆ If you wanted to get the correlation between raters (using Pearson's *rho*) what would u do?

Workbook2

cash

Home Layout Tables Charts SmartArt Formulas Data Review

Font Alignment Number Format Cells Themes

Calibri (Body) 12 Wrap Text General Conditional Formatting Styles Actions Themes Aa Aa

A1 A B C D E F G H I J K L M N O

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30

Sheet1 +

Spreadsheets can be used,  
Pearson's is a known fn, but Cohen's Kappa is not.

# Sentiment Lists

## Q2

# Finding...

- ◆ Do, some searches and find 3 sentiment lists that are commonly used in previous research
- ◆ For 2 of these lists, select 10 positive and 10 negative words (randomly)
- ◆ Evaluate each word, discussing whether it is really positive/negative; for each word try to find sentential context in which it might be interpreted with the opposite valence

# A Program

## Q3

# Practical Q3

- ◆ So, take a look at this sentiment program and see what is happening in the different variables
- ◆ Now consider ways to improve the training.
- ◆ Eg if you removed stop-words from the inputs what do you think might happen
- ◆ Implement this or another solution

```

# Andy Bromberg's Simple Sentiment Analysis System
import re, math, collections, itertools
import nltk, nltk.classify.util, nltk.metrics
from nltk.classify import NaiveBayesClassifier
from nltk.metrics import BigramAssocMeasures
from nltk.probability import FreqDist, ConditionalFreqDist

def evaluate_features(feature_select):
    negSentences = open('/Users/user/Desktop/rt-polarity-neg.txt', 'r')
    posSentences = open('/Users/user/Desktop/rt-polarity-pos.txt', 'r')
    negSentences = re.split(r'\n', negSentences.read())
    posSentences = re.split(r'\n', posSentences.read())

    posFeatures = []
    negFeatures = []
    for i in posSentences:
        posWords = re.findall(r"[\w']+|[.,!?;]", i)
        posWords = [feature_select(posWords), 'pos']
        posFeatures.append(posWords)
    for i in negSentences:
        negWords = re.findall(r"[\w']+|[.,!?;]", i)
        negWords = [feature_select(negWords), 'neg']
        negFeatures.append(negWords)

    posCutoff = int(math.floor(len(posFeatures)*3/4))
    negCutoff = int(math.floor(len(negFeatures)*3/4))
    trainFeatures = posFeatures[:posCutoff] + negFeatures[:negCutoff]
    testFeatures = posFeatures[posCutoff:] + negFeatures[negCutoff:]
    classifier = NaiveBayesClassifier.train(trainFeatures)

    referenceSets = collections.defaultdict(set)
    testSets = collections.defaultdict(set)
    for i, (features, label) in enumerate(testFeatures):
        referenceSets[label].add(i)
        predicted = classifier.classify(features)
        testSets[predicted].add(i)

def make_full_dict(words):
    return dict([(word, True) for word in words])

evaluate_features(make_full_dict)

```

import a bunch  
of stuff

read in +/-  
sentences

turn sents. into  
words and  
structure  
feature inputs  
for classifier

```

# Andy Bromberg's Simple Sentiment Analysis System
import re, math, collections, itertools
import nltk, nltk.classify.util, nltk.metrics
from nltk.classify import NaiveBayesClassifier
from nltk.metrics import BigramAssocMeasures
from nltk.probability import FreqDist, ConditionalFreqDist

def evaluate_features(feature_select):
    negSentences = open('/Users/user/Desktop/rt-polarity-neg.txt', 'r')
    posSentences = open('/Users/user/Desktop/rt-polarity-pos.txt', 'r')
    negSentences = re.split(r'\n', negSentences.read())
    posSentences = re.split(r'\n', posSentences.read())

    posFeatures = []
    negFeatures = []
    for i in posSentences:
        posWords = re.findall(r"[\w']+|[.,!?;]", i)
        posWords = [feature_select(posWords), 'pos']
        posFeatures.append(posWords)
    for i in negSentences:
        negWords = re.findall(r"[\w']+|[.,!?;]", i)
        negWords = [feature_select(negWords), 'neg']
        negFeatures.append(negWords)

    posCutoff = int(math.floor(len(posFeatures)*3/4))
    negCutoff = int(math.floor(len(negFeatures)*3/4))
    trainFeatures = posFeatures[:posCutoff] + negFeatures[:negCutoff]
    testFeatures = posFeatures[posCutoff:] + negFeatures[negCutoff:]
    classifier = NaiveBayesClassifier.train(trainFeatures)

    referenceSets = collections.defaultdict(set)
    testSets = collections.defaultdict(set)
    for i, (features, label) in enumerate(testFeatures):
        referenceSets[label].add(i)
        predicted = classifier.classify(features)
        testSets[predicted].add(i)

def make_full_dict(words):
    return dict([(word, True) for word in words])

evaluate_features(make_full_dict)

```

make 2/3s of  
features  
training set

do training

run the test set  
on the trained  
system

```

def evaluate_features(feature_select):
    negSentences = open('/Users/user/Desktop/rt-polarity-neg.txt', 'r')
    posSentences = open('/Users/user/Desktop/rt-polarity-pos.txt', 'r')
    negSentences = re.split(r'\n', negSentences.read())
    posSentences = re.split(r'\n', posSentences.read())

    posFeatures = []
    negFeatures = []
    for i in posSentences:
        posWords = re.findall(r"[\w']+|[.,!?;]", i)
        posWords = [feature_select(posWords), 'pos']
        posFeatures.append(posWords)
    for i in negSentences:
        negWords = re.findall(r"[\w']+|[.,!?;]", i)
        negWords = [feature_select(negWords), 'neg']
        negFeatures.append(negWords)

    posCutoff = int(math.floor(len(posFeatures)*3/4))
    negCutoff = int(math.floor(len(negFeatures)*3/4))
    trainFeatures = posFeatures[:posCutoff] + negFeatures[:negCutoff]
    testFeatures = posFeatures[posCutoff:] + negFeatures[negCutoff:]
    classifier = NaiveBayesClassifier.train(trainFeatures)

    referenceSets = collections.defaultdict(set)
    testSets = collections.defaultdict(set)
    for i, (features, label) in enumerate(testFeatures):
        referenceSets[label].add(i)
        predicted = classifier.classify(features)
        testSets[predicted].add(i)

def make_full_dict(words):
    return dict([(word, True) for word in words])

evaluate_features(make_full_dict)

```

create a dict  
structure for  
the words in  
sentence

call make\_full\_dict  
as the arg to  
evaluate\_features

# Other Bits for Outputs

#Outputs

```
print('train on %s instances, test on %s instances' % (len(trainFeatures), len(testFeatures)))
print('accuracy:', nltk.classify.util.accuracy(classifier, testFeatures))
print('pos precision:', nltk.metrics.precision(referenceSets['pos'], testSets['pos']))
print('pos recall:', nltk.metrics.recall(referenceSets['pos'], testSets['pos']))
print('neg precision:', nltk.metrics.precision(referenceSets['neg'], testSets['neg']))
print('neg recall:', nltk.metrics.recall(referenceSets['neg'], testSets['neg']))
classifier.show_most_informative_features(10)
```

Python 3.4.1 Shell

```
Python 3.4.1 (default, May 21 2014, 01:39:38)
[GCC 4.2.1 Compatible Apple LLVM 5.1 (clang-503.0.40)] on darwin
Type "copyright", "credits" or "license()" for more information.
>>> ===== RESTART =====
>>>
using all words as features
train on 7998 instances, test on 2666 instances
accuracy: 0.77344336084021
pos precision: 0.7881422924901186
pos recall: 0.7479369842460615
neg precision: 0.7601713062098501
neg recall: 0.7989497374343586
Most Informative Features
engrossing = True          pos : neg    =      17.0 : 1.0
quiet = True                pos : neg    =      15.7 : 1.0
mediocre = True             neg : pos    =      13.7 : 1.0
absorbing = True            pos : neg    =      13.0 : 1.0
portrait = True              pos : neg    =      12.4 : 1.0
flaws = True                pos : neg    =      12.3 : 1.0
refreshing = True            pos : neg    =      12.3 : 1.0
inventive = True             pos : neg    =      12.3 : 1.0
refreshingly = True          pos : neg    =      11.7 : 1.0
triumph = True               pos : neg    =      11.7 : 1.0
>>> |
```

```

import re, math, collections, itertools
import nltk, nltk.classify.util, nltk.metrics
from nltk.classify import NaiveBayesClassifier
from nltk.metrics import BigramAssocMeasures
from nltk.probability import FreqDist, ConditionalFreqDist

def evaluate_features(feature_select):
    #reading pre-labeled input and splitting into lines
    negSentences = open('/Users/user/Desktop/rt-polarity-neg.txt', 'r')
    posSentences = open('/Users/user/Desktop/rt-polarity-pos.txt', 'r')
    negSentences = re.split(r'\n', negSentences.read())
    posSentences = re.split(r'\n', posSentences.read())

    posFeatures = []
    negFeatures = []
    # breaks up the sentences into lists of individual words
    # creates instance structures for classifier
    for i in posSentences:
        posWords = re.findall(r"[\w']+|[.,!?;]", i)
        posWords = [feature_select(posWords), 'pos']
        posFeatures.append(posWords)
    for i in negSentences:
        negWords = re.findall(r"[\w']+|[.,!?;]", i)
        negWords = [feature_select(negWords), 'neg']
        negFeatures.append(negWords)

    posCutoff = int(math.floor(len(posFeatures)*3/4))
    negCutoff = int(math.floor(len(negFeatures)*3/4))
    trainFeatures = posFeatures[:posCutoff] + negFeatures[:negCutoff]
    testFeatures = posFeatures[posCutoff:] + negFeatures[negCutoff:]

    #Runs the classifier on the testFeatures
    classifier = NaiveBayesClassifier.train(trainFeatures)

    #Sets up labels to look at output
    referenceSets = collections.defaultdict(set)
    testSets = collections.defaultdict(set)
    for i, (features, label) in enumerate(testFeatures):
        referenceSets[label].add(i)
        predicted = classifier.classify(features)
        testSets[predicted].add(i)

    #Outputs
    print('train on %s instances, test on %s instances' % (len(trainFeatures), len(testFeatures)))
    print('accuracy:', nltk.classify.util.accuracy(classifier, testFeatures))
    print('pos precision:', nltk.metrics.precision(referenceSets['pos'], testSets['pos']))
    print('pos recall:', nltk.metrics.recall(referenceSets['pos'], testSets['pos']))
    print('neg precision:', nltk.metrics.precision(referenceSets['neg'], testSets['neg']))
    print('neg recall:', nltk.metrics.recall(referenceSets['neg'], testSets['neg']))
    classifier.show_most_informative_features(10)

def make_full_dict(words):
    return dict([(word, True) for word in words])

print('using all words as features')
evaluate_features(make_full_dict)

```

# Practical Q3

- ◆ So, take a look at the program and see what is happening in the different variables
- ◆ Now consider ways to improve the training.
- ◆ Eg if you removed stop-words from the inputs what do you think might happen
- ◆ Implement this or another solution